

Communications and Computer Networks

The fundamental purpose of a communication system is the exchange of data between two parties. An example is the exchange of voice signals between two telephones on the same network. It is often impractical for two communicating devices to be directly, point-to-point connected. This is for one (or both) of the following contingencies:

1. The devices are very far apart, and hence cannot share a dedicated link.
2. There is a set of devices, each of which requires a link to many of the others at various times.

The solution to this problem is to attach each device to a communication network. The two categories into which communication networks are classified: local area networks (LANs) and wide area networks (WANs).

Wide Area Networks

Wide area networks generally cover a large geographical area, typically multiple cities, countries, or continents. Typically, a WAN consists of a number of interconnected switching nodes called *routers*. A transmission from any device is routed through these internal nodes to the specified destination device. These switching nodes are not concerned with the content of the data; rather, their purpose is to provide a switching facility that will move the data from node to node until they reach their destination. The switching employed is usually:

1. **Circuit Switching:** In a circuit switching network, a dedicated communications path is established between two stations through the nodes of the network. Data generated by the source station are transmitted along the dedicated path as rapidly as possible. Since the path is already established *a priori* there is no delay involved. The most common example of circuit switching is the telephone network. Disadvantages are the cost involved; for example one pays a fixed rate for a phone call even when the two parties do not talk.
2. **Packet Switching:** Data is sent out in a sequence of small chunks called *packets*. Each packet is passed through the network from node to node along some path leading from source to destination. At each switching

node, the entire packet is received, stored briefly in a queue, and then transmitted to the next node. This is commonly used in computer to computer communications. An example of packet switching is the postal network. Whenever the network becomes overloaded, computers using the network must wait before they can send additional packets. However since multiple computers can share the network bandwidth, fewer connections are required and the cost is kept low.

Usually, WANs range from 1.5 Mbps to 155 Mbps (where Mbps is million bits per second). Examples of WAN technologies are: ARPANET (Section 2.7 of Comer [1]), X.25 (Section 2.11.1 of Comer [1]), Frame Relay, ISDN (Integrated Services Digital Network), ATM (Asynchronous Transfer Mode) (Section 2.6 of Comer [1]).

Local Area Networks

The scope of the LAN is small, typically a single building or a cluster of buildings. It is also usually the case that the LAN is owned by the same organization that owns the communicating computers. The difference in geographic scope of LANs leads to technologies that are different from those of WANs.

Traditionally, LANs make use of a broadcast network approach rather than a switching network. With a broadcast communication network, there are no intermediate switching nodes. There is a Network Interface Card (NIC) that connects a computer directly to the network (medium) shared by other stations. A transmission from any one station is broadcast to and received by all other stations. Data is usually transmitted in packets. Because the medium is shared, only one station can transmit a packet at a time.

The internal data rates for LANs are typically much greater than those of WANs. Typically, this number is 10 Mbps to 2 Gbps (where Gbps is billion bits per second).

Examples of LAN technologies include: Ethernet (IEEE 802.3) (Section 2.4 of Comer [1]), FDDI (token ring passing LAN) (Section 2.5 of Comer [1]).

Internetworking concept, protocols and architecture

The primary goal of internetworking is a system that hides the details of the underlying network hardware while providing universal communication services.

There are two approaches to hiding network details.

1. **Application-Level Interconnection:** In such systems, an application program executing on each computer in the network understands the details of the network connections, and interoperates across those connections with application programs on other computers. Take for instance an electronic mail system, where an application mail program is configured to forward a message to another mail program on another computer. The path from source to destination may involve many different networks, but this does not matter as long as the mail systems on the various computers cooperate by forwarding the message. This, however, leads to cumbersome communication for the following reasons:

- Adding new functionality to the system means building a new application program for each computer.
- Adding new network hardware means modifying existing application programs on the computer.
- Finally, each application program on the computer would need to understand the network connections for the computer, resulting in duplication of code.

2. **Network-Level Interconnection:** A network-level interconnection provides a mechanism that delivers small packets of data from the original source to destination without using intermediate application programs. This separates the data communication duties from application programs, and existing network technologies can be modified, while the application programs on computers remain unchanged.

The key to designing universal network-level interconnection is the concept of internetworking. The idea is to detach the notions of communication from the details of network technologies and hide low-level details from the user. Thus, instead of a single module for performing communications, there is a *structured layer* of modules that implements the communication function. Each module/layer in a computer communicates with its peer module/layer on another computer via a set of *protocols*, which may be defined as set of rules governing the exchange of data between two entities. The layers and protocols together define the *network architecture*.

For example an English philosopher and a French philosopher wish to communicate, but neither speaks the other's language. Furthermore, neither has on hand a translator that can translate to the language of the other. However, both philosophers have Dutch translators on their staffs. Each translator has a secretary, who communicates with other translator's secretary via fax. Let us suppose the English translator has a message for his French counterpart, and his message in English reads *I like rabbits*. The philosopher conveys this message to his translator, who translates the message into Dutch which reads *Ik Hou Van Konijnen*. The translator also adds a header to this message for his counterpart, which indicates that the language employed is Dutch. The translator gives his information to the secretary who adds the fax number to the message and faxes this information to her counterpart. The other secretary on receipt of this message, removes the fax number from the message, and passes on the remaining information to the French translator who, in turn, translates the Dutch message into an equivalent French message that reads *J'aime les lapins*. He passes this message to the French philosopher. Here, both the source and the destination have three layers of communication. Layers 1, 2 and 3 represent the secretaries, translators, and philosophers respectively. The protocol employed between the peer to peer translator layers is the Dutch language, while fax is the protocol employed by the secretary layers. Now suppose that no Dutch translators are available, but each philosopher has instead Hindi translators on their staffs. The communication between the two philosophers will still take place, and the only change is that layer 2 represents Hindi translators with the peer to peer protocol being the Hindi language. Furthermore, this change does not affect the philosopher and secretary layers.

In very general terms communications involve three agents: applications, computers, and networks. One example of an application is a file transfer (ftp)

operation. These applications execute on computers, that can often support multiple simultaneous applications. Computers, in turn, are connected to networks, and the data to be exchanged are transferred by the network from one computer to another. With these three agents in mind, it is natural to organize the communication task into three relatively independent layers.

1. **Network Access layer:** This layer is concerned with the exchange of data between a computer and the network to which it is attached. The sending computer must provide the network with the address of the destination computer, so that the network will route the data to the appropriate destination. The sending computer may also wish to invoke certain services, such as priority, that might be provided by the network. The specific software employed at this layer depends on the type of network used, and different standards developed for circuit/packet switching etc. Thus, the layers above the network access layer need not be concerned about the specifics of the network to be used. Thus, two computers attached to different physical networks can now communicate.
2. **Transport layer:** Regardless of the nature of the applications that are exchanging data (ftp, elm), there is usually a requirement that data be exchanged reliably. Thus, it makes sense to collect those mechanisms in a common layer called the transport layer.
3. **Application layer:** The application layer contains the logic needed to support various user applications. For each different type of application, a separate module is needed that is characteristic to that application.

Moreover, these three layers communicate with their peers on another computer via a set of protocols analogous to the communication between the two philosophers mentioned above. These protocols require the exchange of various control information, including the network address of the destination computer, and an address that is specific to the particular application (The latter addresses are called service access points (SAPs) or ports). Suppose an application, associated with SAP 1 on computer X wishes to send a message to another application associated with SAP 2 on computer Y. The application at X hands the message to its transport layer. The transport layer may break this block into smaller pieces to make it more manageable. To each of these pieces, the transport layer appends a transport header containing protocol control information. This process is known as *encapsulation*, and the combination of data from the next higher layer and control information is known as a protocol data unit (PDU); in this case it is referred to as a transport PDU. For example, the header may contain the following information:

1. **Destination SAP:** When the destination transport layer receives the transport PDU it must know to whom the data is to be delivered.
2. **Sequence Number:** Because the transport protocol is sending a sequence of PDUs, it numbers them sequentially so that if they arrive out of order, the destination transport layer may reorder them.
3. **Error-detection code:** The sending transport layer may also include a code that is a function of the contents of the data portion of the PDU. The receiving transport layer protocol performs the same calculation. If there is a discrepancy, the receiver can discard the PDU and take corrective action.

The transport layer hands each PDU to the network access layer. This layer appends further control information to the transport PDU (which is regarded as the data portion) creating a network access PDU. For example, the header may contain the following information:

1. **Destination computer address:** The network must know to which computer the data must be delivered.
2. **Facilities requests:** The network access protocol might want to make use of certain network facilities, such as priority.

The network access layer is not concerned with the details of the transport header, such as the destination SAP address for instance. The network accepts the network PDU from computer X and delivers it to computer Y. The network access module in Y receives the PDU, strips off the header, and transfers the enclosed transport PDU to the transport layer module. The transport layer examines the transport PDU header, and, on the basis of the SAP field delivers the enclosed record to the appropriate application, in this case the file transfer module in computer Y.

Common protocol architectures

Two protocol architectures have served the basis for the development of communication standards:

1. The TCP/IP protocol suite.
2. The OSI reference model.

We provide a brief overview of these two architectures in this section. Finally, we highlight the important differences between these two models.

The TCP/IP protocol suite

TCP/IP is the result of protocol research and development conducted on the experimental packet-switched network, ARPANET, funded by the Defense Advanced Research Projects Agency (DARPA), and is generally referred to the TCP/IP protocol suite.

The communication task in TCP/IP is organized into five relatively independent layers:

1. **Physical Layer:** The physical layer is the network hardware layer, and is concerned with the characteristics of the transmission medium, the nature of the signals, data rate, and related matters.
2. **Network Interface Layer:** This is the lowest software layer, responsible for accepting IP datagrams, and transmitting them over a specific network in frames.
3. **Internet Layer:** The internet layer handles communication from one machine to another. It accepts a request to send a packet from the transport layer along with an identification of the machine to which the packet is to be sent. It encapsulates the packet in an IP datagram, uses a routing algorithm to determine whether to deliver the datagram directly or send

it to a router, and passes the datagram to the network interface layer for transmission. It also receives datagrams, checks their validity, and uses a routing algorithm to decide whether the datagram is to be processed locally or sent to another router. For datagrams addressed to the local machine, the internet layer deletes the datagram header and delivers the packet information to the appropriate transport layer. The internet layer adopts the internet protocol (IP) (the IP in TCP/IP). The aim of the internet layer is to provide *best-effort delivery*. It does not attempt to correct errors, although it sends ICMP (Internet Control Message Protocols) error and control messages as needed.

4. **Transport Layer:** It's primary duty is to provide communication between the application program on the source computer with that on the destination computer. It divides the stream of data being transmitted by the application program into small packets, and passes each packet along with a destination computer address to the internet layer for transmission. It also provides reliable transport at the receiving end ensuring that the data arrives without error and in sequence.
5. **Application Layer:** This contains the necessary logic for a particular application. The application layer interacts with the transport layer to send and receive data.

The application and transport layers are called *end-to-end*, since they are only implemented on the source and the destination computers. A router, for instance, need not have a transport and an application layer. Typically, the transport, internet and network interface layers are implemented in software, and in the operating system of the computer. The application layer is a software layer that resides in the user disc space.

Common examples of TCP/IP protocols include:

1. **Application Layer:** Telnet, FTP, e-mail etc.
2. **Transport Layer:** TCP (Transmission Control Protocol), UDP (User Datagram Protocol).
3. **Internet Layer:** IP (Internet Protocol), ICMP (Internet Control Message Protocol), IGMP (Internet Group Management Protocol).
4. **Network Interface Layer:** Device driver and interface card.

The OSI reference model

The OSI (Open Systems Interconnection) proposed a standard known as the OSI reference model. It contains 7 conceptual layers organized as follows:

1. Physical Hardware Connection.
2. Data Link Layer.
3. Network Layer.
4. Transport Layer.
5. Session Layer.

6. Presentation Layer.

7. Application Layer.

The OSI reference model was designed as a framework for developing protocol standards. The designers of OSI assumed that this model and the protocols developed within this model would dominate computer communications, eventually replacing rival multivendor models such as TCP/IP. However this has not happened. It is the TCP/IP architecture that has come to dominate. Thus, our emphasis in this course will be on TCP/IP. The book by Tanenbaum [3] is a good introduction to the OSI 7 layer model.

Differences between TCP/IP and OSI models

There are two subtle and important differences between the TCP/IP layering and the OSI reference model. This section will provide a brief overview of the essential differences; more details can be found in Section 11.6 of Comer [1].

1. **Link-Level vs. End-To-End Reliability:** In the OSI model protocol software detects and handles errors at all layers. In contrast to such a scheme, TCP/IP assumes that reliability is an end-to-end problem. There is little or no reliability in most TCP/IP network interface software; rather this is the duty of the transport layer. The resulting freedom from interface layer verification makes TCP/IP software much easier to understand and implement correctly.
2. **Locus of Intelligence and Decision Making:** The OSI layer model assumes that the network is a utility that provides transport service. The network vendor also handles problems like routing, flow control, and acknowledgments, internally, making transfers reliable. In short, the network is a complex independent system to which one may attach relatively simple computers; the hosts themselves participate minimally in the network operation. In contrast, TCP/IP requires hosts to participate in almost all the network protocols, such as end-to-end error detection and recovery. Thus, a TCP/IP internet can be viewed as a simple packet delivery system to which intelligent hosts attach.

Suggested Readings

1. Chapters 2,3 and 11 of Comer [1].
2. Chapters 1 and 2 of Stallings [2].
3. Chapter 1 of Tanenbaum [3].

References

- [1] D.E. COMER, *Internetworking with TCP/IP: Principles, Protocols, and Architectures*, 4th edition, Prentice Hall, NJ, 2000.
- [2] W. STALLINGS, , *Data & Computer Communications*, 6th edition, Prentice Hall, NJ, 2002.
- [3] A. TANENBAUM, *Computer Networks*, Prentice Hall, NJ, 2000.